

Deep Reinforcement Learning: An Overview

Patrick Emami

PhD student, CISE department

February 24, 2018

Motivation

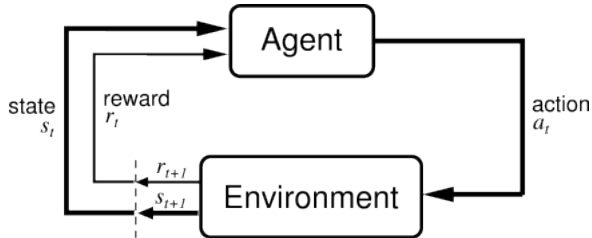
What is a good framework for studying intelligence?

Motivation

What is a good framework for studying intelligence?

What are the necessary and sufficient ingredients for building agents that learn and act like people?

Reinforcement Learning

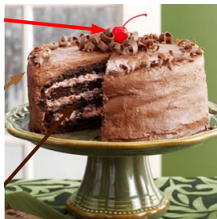


Source: Sutton, Richard S., and Andrew G. Barto. Reinforcement learning: An introduction. Vol. 1. No. 1. Cambridge: MIT press, 1998.

Reinforcement Learning

My Perspective

Reinforcement Learning is necessary but not sufficient for general (strong) artificial intelligence



Markov Decision Processes

Definition

A **Markov decision process** (MDP) is a formal way to describe the sequential decision-making problems encountered in RL.

Markov Decision Processes

Definition

A **Markov decision process** (MDP) is a formal way to describe the sequential decision-making problems encountered in RL.

In the simplest RL setting, an MDP is specified by states S , actions A , an episode length H , and a reward function $r(s, a)$.

Policies and Value Functions

- A policy $\pi(a|s)$ is a behavior function for selecting an action given the current state.
- The **action-value function** is the expected total reward accumulated from starting in state s , taking action a , and following policy π until the end of the length H episode:

$$Q^\pi(s, a) = \mathbb{E}_\pi \left[\sum_{t=0}^H r_t | s_t = s, a_t = a \right]$$

“What is the utility of doing action a when I’m in state s ?”

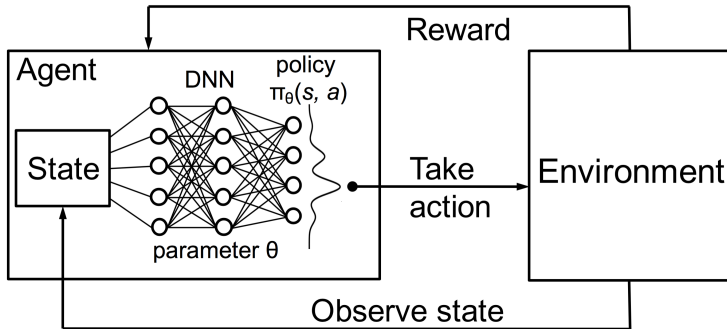
Big Picture

Find policy π^* that maximizes expected total reward, i.e.,

$$\pi^* = \operatorname{argmax}_{\pi} Q^{\pi}(s, a).$$

In particular, for any start state $s_0 \in S$, the agent can use π^* to select the action a_0 that will maximize its expected total reward.

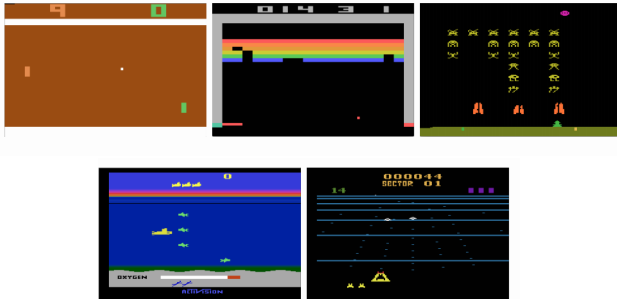
Deep Reinforcement Learning



Deep Reinforcement Learning



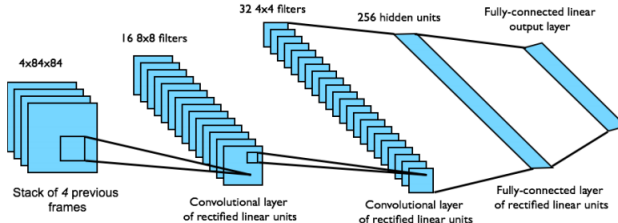
Playing Atari (2013)



Source: Mnih, Volodymyr, et al. "Playing atari with deep reinforcement learning." arXiv preprint arXiv:1312.5602 (2013).

Deep Q-Network (DQN)

- ▶ End-to-end learning of values $Q(s, a)$ from pixels s
- ▶ Input state s is stack of raw pixels from last 4 frames
- ▶ Output is $Q(s, a)$ for 18 joystick/button positions
- ▶ Reward is change in score for that step



Network architecture and hyperparameters fixed across all games
[Mnih et al.]

Just Apply Gradient Descent

- Represent $Q^\pi(s, a)$ by a **deep Q-network** with weights w

$$Q(s, a, w) \approx Q^\pi(s, a)$$

- Define objective function by mean-squared **Bellman error**

$$L(w) = \mathbb{E} \left[\left(r + \gamma \max_{a'} Q(s', a', w) - Q(s, a, w) \right)^2 \right]$$

- Leading to the following gradient

$$\frac{\partial L}{\partial w} = \mathbb{E} \left[\left(r + \gamma \max_{a'} Q(s', a', w) - Q(s, a, w) \right) \frac{\partial Q(s, a, w)}{\partial w} \right]$$

- Optimize with stochastic gradient descent

Stability Issues with Deep RL

Naive Q-learning with non-linear function approximation **oscillates** or **diverges**

- Experiences from episodes generated during training are correlated, non-iid
- Policy can change rapidly with slight changes to Q-values
- Q-learning gradients can be large and unstable when backpropagated

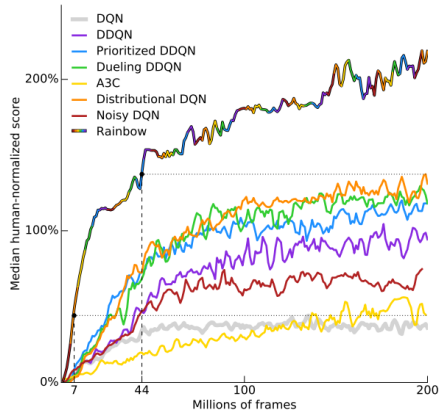
Stabilizing Deep RL

- Maintain a **replay buffer** of experiences to uniformly sample from to compute gradients for Q-network. This decorrelates samples and improves samples efficiency
- Hold the parameters of the target Q-values fixed in Bellman error with a **target Q-network**. Periodically update the parameters of the target network
- Clip rewards and potentially clip gradients as well

	B. Rider	Breakout	Enduro	Pong	Q*bert	Seaquest	S. Invaders
Random	354	1.2	0	-20.4	157	110	179
Sarsa [3]	996	5.2	129	-19	614	665	271
Contingency [4]	1743	6	159	-17	960	723	268
DQN	4092	168	470	20	1952	1705	581
Human	7456	31	368	-3	18900	28010	3690
HNeat Best [8]	3616	52	106	19	1800	920	1720
HNeat Pixel [8]	1332	4	91	-16	1325	800	1145
DQN Best	5184	225	661	21	4500	1740	1075

Source: Mnih, Volodymyr, et al. "Playing
 atari with deep reinforcement learning."
 arXiv preprint arXiv:1312.5602 (2013).

Rainbow DQN (2017)



Source: Hessel, Matteo, et al. "Rainbow: Combining Improvements in Deep Reinforcement Learning." arXiv:1710.02298 (2017).

Figure 1: Median human-normalized performance across 57 Atari games. We compare our integrated agent (rainbow-

AlphaGo (2015)

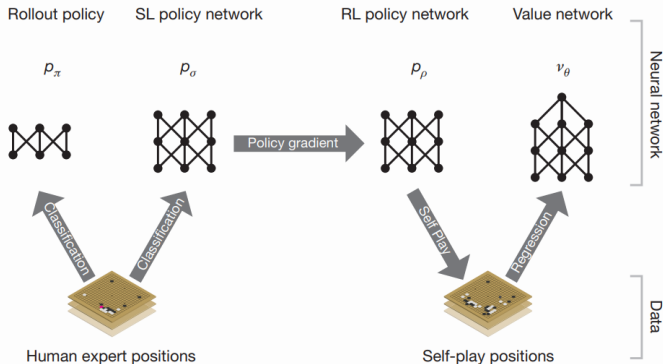


It was thought that AI was a decade away from beating humans at Go

AlphaGo

Key Ingredients

Tree search augmented with policy and value deep networks that intelligently control exploration and exploitation



Monte Carlo Tree Search

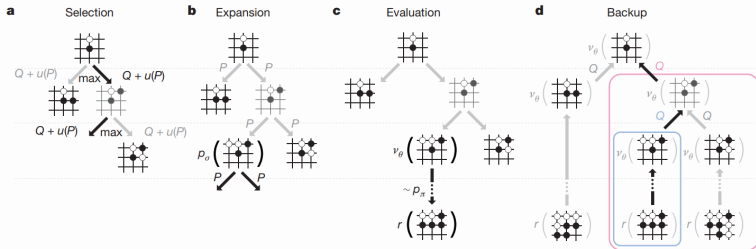


Figure 3 | Monte Carlo tree search in AlphaGo. **a**, Each simulation traverses the tree by selecting the edge with maximum action value Q , plus a bonus $u(P)$ that depends on a stored prior probability P for that edge. **b**, The leaf node may be expanded; the new node is processed once by the policy network p_θ , and the output probabilities are stored as prior probabilities P for each action. **c**, At the end of a simulation, the leaf node

is evaluated in two ways: using the value network v_θ ; and by running a rollout to the end of the game with the fast rollout policy p_π , then computing the winner with function r . **d**, Action values Q are updated to track the mean value of all evaluations $r(\cdot)$ and $v_\theta(\cdot)$ in the subtree below that action.

Source: Silver, David, et al. "Mastering the game of Go with deep neural networks and tree search." *nature* 529.7587 (2016): 484-489.

Continuous control

① Locomotion Behaviors

<https://www.youtube.com/watch?v=g59nSURxYgk>

② Learning to Run

https://www.youtube.com/watch?v=mbjuaRg__DI

③ Robotics

<https://www.youtube.com/watch?v=Q4bMcUk6pcw&t=56s>

Continuous control

Can we learn Q^{π^*} by minimizing the expected Bellman error?

Continuous Action Spaces

For $A \in \mathbb{R}^n$,

$$\mathbb{E} \left[\left(r + \gamma \max_{a' \in A} Q(s', a', w) - Q(s, a, w) \right)^2 \right]$$

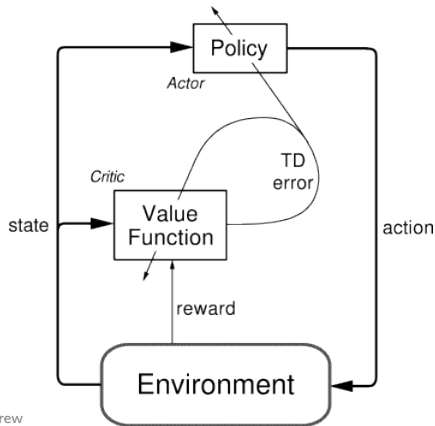
Requires solving a non-convex optimization problem!

Deep Deterministic Policy Gradient (2016)

- 1 Let the policy be a deterministic function $\pi(s, \theta) : S \rightarrow A$, $S \in \mathbb{R}^m$, $A \in \mathbb{R}^n$, parameterized as a deep network
- 2 Still maximize expected total reward, except now need to compute the deterministic policy (actor) gradient and the (critic) action-value gradient
- 3 Train both the policy and action-value networks with an **actor-critic** approach

Source: Lillicrap, Timothy P., et al.
"Continuous control with deep
reinforcement learning." arXiv preprint
arXiv:1509.02971 (2015).

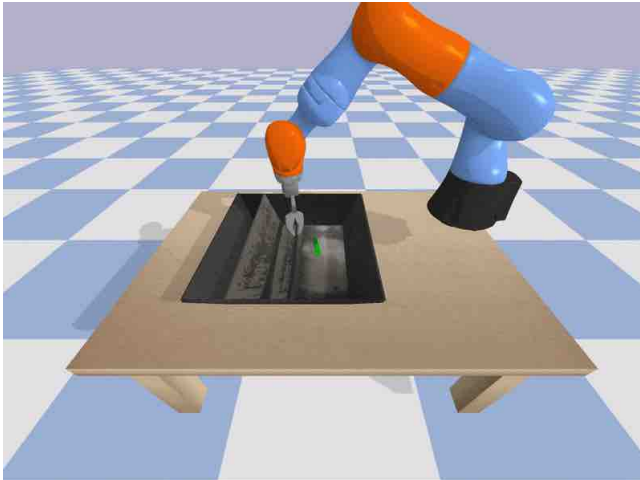
Actor-Critic



Source: Sutton, Richard S., and Andrew G. Barto. Reinforcement learning: An introduction. Vol. 1. No. 1. Cambridge: MIT press, 1998.

Future Research Directions

- ① Sample-efficient learning by embedding priors about the world, e.g., intuitive physics
- ② Low-variance, unbiased policy gradient estimators
- ③ Multi-agent RL
- ④ Safe RL
- ⑤ Reinforcement learning on combinatorial action spaces



Source:
<http://blog.otoro.net/2017/11/12/evolving-stable-strategies/>

Fin.